
DeadlineIQ

Cognitive-Load AI Task Planner & Focus Companion

Technical Documentation

Version 1.0 • June 2026

Author: Moteesh A

AI-Powered

Firebase

Gemini 2.5

React 18

Contents

1	Product Overview	3
1.1	Introduction	3
1.2	Key Capabilities	3
2	User Interface Walkthrough	5
2.1	Primary Navigation HUD (Dashboard)	5
2.2	Local Neural Network Risk Model	6
2.3	Capacity Mapping & Focus Nudges	6
2.4	Gemini Natural Language Task Parser	7
2.5	Interactive Co-pilot IQ Coach Sidebar	7
2.6	Task Deconstruction & Subtask Builder	8
2.7	Kanban Progress Board	8
2.8	Glassmorphic Calendar Grid	8
2.9	Procrastination Forensics & Score Forecasting	9
2.10	Habit Streak Tracker	9
2.11	Chrome Extension Panel	10
3	Google Technologies Integration	11
3.1	Integration Summary	11
3.2	Architecture Dependency Diagram	12
4	Technical Architecture	13
4.1	Front-End Stack	13
4.2	Client-Side Neural Network (MLP)	13
4.2.1	Model Specification	14
4.3	Firestore Security Rules	14
4.4	Project Directory Structure	15
5	Local Development Setup	16
5.1	Prerequisites	16
5.2	Installation Steps	16
5.2.1	Step 1 Navigate and Install Dependencies	16
5.2.2	Step 2 Environment Configuration	16
5.2.3	Step 3 Start the Development Server	17

6	Deployment	18
6.1	Docker Deployment	18
6.1.1	Build the Container Image	18
6.1.2	Run the Container	18
6.2	Firebase Hosting Deployment	18
6.3	Google Cloud Run Deployment	19
7	Security Considerations	20
7.1	Authentication Layer	20
7.2	Data Isolation at the Database Layer	20
7.3	Client-Side ML Privacy	20
7.4	API Key Security Best Practices	20

Product Overview

1.1 Introduction

DeadlineIQ is an ultra-premium, AI-driven productivity companion engineered to defeat procrastination and optimise cognitive focus. Unlike conventional productivity tools that rely on passive push alerts which users habitually ignore, DeadlineIQ introduces a *proactive, spatial, and cognitively aware* approach to task management.

The system analyses behavioural patterns in real time, predicts task-completion risks using a client-side machine learning model, and orchestrates personalised focus windows directly around the user's existing calendar commitments.

DeadlineIQ replaces reactive reminders with **predictive intelligence** – it intervenes before deadlines are missed, not after. Every feature is designed to reduce cognitive load, not add to it.

1.2 Key Capabilities

- 1. Local Neural Network Risk Scoring** A client-side Multi-Layer Perceptron (MLP) computes personalised procrastination probability scores without transmitting sensitive data to external servers.
- 2. Natural Language Task Parsing** Powered by *Gemini 2.5 Flash*, users dictate chaotic schedules in plain English and receive structured task objects with cognitive-load estimates and confidence scores.
- 3. Calendar-Aware Focus Scheduling** Deep integration with the Google Calendar API allocates focus windows around existing commitments and peak-productivity periods.
- 4. Procrastination Forensics** Recalculates behavioural delays to identify and categorise procrastination root causes (e.g., *Fear of Failure*, *Task Ambiguity*) with SVG trend charts.

- 5. Agentic Co-pilot (IQ Coach)** A sliding sidebar AI agent executes scheduling actions, reschedules tasks, and commits all updates to Firestore in real time.
- 6. Habit Streak Tracking** Glowing circular gauges monitor routine completions with Gemini-powered success probability forecasts.
- 7. Chrome Extension Integration** Synchronises browser research directly to the cloud task dashboard without context-switching.

CHAPTER 2

User Interface Walkthrough

DeadlineIQ's UI is built on a **VisionOS-inspired spatial design language** floating telemetry modules, glassmorphic panels, and a premium dark-mode slate palette that minimises distraction while maximising information density.

2.1 Primary Navigation HUD (Dashboard)

The main dashboard is the command centre of DeadlineIQ. It surfaces live task metrics, multi-agent status, and focus-window data through floating telemetry panels.




Figure 2.1: Primary HUD showing annotated key regions: Active AI banner (top) and floating controls (bottom right).

Key UI Elements:

- Floating telemetry cards with live task counts and urgency indicators
- Multi-agent status banner showing active AI operations
- Focus score ring and daily priority hour counter
- One-tap access to all feature modules via the sidebar rail

2.2 Local Neural Network Risk Model

An integrated client-side neural net assesses active task parameters in real time to produce a personalised **Procrastination Probability Score**.




Figure 2.2: Local MLP Neural Network widget showing the risk classification highlight.

2.3 Capacity Mapping & Focus Nudges

This module displays floating telemetry panels indicating current focus scores, active priority hours, and ambient audio controls that help users enter deep-work states.




Figure 2.3: Capacity Dial & Audio Telemetry layout showing workload limits.

2.4 Gemini Natural Language Task Parser

An automated schedule deconstruction module. Gemini 2.5 Flash parses the voice or text input, extracts structured task objects, and estimates cognitive load metrics and confidence scores automatically.




Figure 2.4: Gemini NL Task Parser showing text-to-task deconstruction and semantic structure evaluation.

User input: *“Finish presentation slides by tomorrow at 4 PM, high priority, will take 3 hours.”*

Gemini output:

- title: Finish presentation slides
- deadline: Tomorrow 16:00
- priority: High
- estimated_duration: 3h
- cognitive_load: 7.4/10
- confidence: 94%

2.5 Interactive Co-pilot IQ Coach Sidebar

A sliding AI sidebar where users can naturally command, reschedule, and complete tasks. The agent executes all actions and commits updates to Firestore in real time.






Figure 2.5: Speech integration and co-pilot activation buttons shown at original high-definition scale.

2.6 Task Deconstruction & Subtask Builder

Gemini automatically decomposes complex deadlines into discrete actionable subtasks. Users can edit subtask parameters, adjust time estimates, and track incremental completion.



(a) Task Deconstruction Modality



(b) Subtask Parameter Builder

Figure 2.6: Task Deconstruction and Subtask Builder interfaces

2.7 Kanban Progress Board

A drag-and-drop board for tracking commitments across progress columns. Audio chimes fire on task completion to provide positive reinforcement.




Figure 2.7: Kanban Progress Board with columns: *Due Today*, *In Progress*, *Completed*

2.8 Glassmorphic Calendar Grid

An edge-to-edge borderless weekly calendar that overlays tasks alongside Google Calendar events and highlights Peak Focus Windows for optimal scheduling.




Figure 2.8: Glassmorphic VisionOS Calendar Grid with Google Calendar sync

2.9 Procrastination Forensics & Score Forecasting

This analytics module catalogues behavioural delay patterns to identify the user's *procrastination fingerprint* and charts predictive score trends with SVG visualisations.



Figure 2.9: Procrastination Forensics and Predictive Score Dashboard

2.10 Habit Streak Tracker

Monitors routine completions with glowing circular progress gauges and integrates Gemini-powered success probability forecasts for each tracked habit.




Figure 2.10: Habit Streak Tracker with circular progress gauges

2.11 Chrome Extension Panel

The companion Chrome extension synchronises web research tabs, captures snippets, and pushes tasks directly to the cloud dashboard eliminating context-switching.




Figure 2.11: Chrome Extension Panel browser-to-dashboard task pipeline

Google Technologies Integration


DeadlineIQ integrates **seven core Google technologies** spanning Cloud, Firebase, and Workspace products to power its intelligence and cloud infrastructure.

3.1 Integration Summary

Technology		Role in DeadlineIQ
Firebase Authentication		Google OAuth 2.0 identity management and session token validation for all protected operations.
Cloud Firestore		Real-time NoSQL database for user profiles, task documents, calendar slots, and behavioural metrics.
Gemini (gemini-2.5-flash) API		Drives NL task parsing, automatic subtask generation, and cognitive-load forecasting.
Google Calendar API		Syncs allocated focus slots and task deadlines to the user's primary Google Calendar.
Google Sheets API		Exports productivity logs and completed task data as Sheets-compatible CSV with OAuth sync validation.
Google Cloud Run		Pre-configured containerised deployments for serverless, auto-scaling production hosting.
Firebase Hosting		Global CDN delivery of production assets under verified SSL (*.web.app / *.firebaseapp.com).

Table 3.1: Google technology stack roles and responsibilities

3.2 Architecture Dependency Diagram



Technical Architecture

4.1 Front-End Stack

Layer	Technology	Notes
UI Framework	React 18	Component-based SPA with concurrent rendering
Build Tool	Vite	Sub-second HMR; optimised production bundles
Styling	Tailwind CSS	Glassmorphism utilities + custom spatial transitions
Client ML	Pure JavaScript	MLP in <code>src/utils/localML.js</code> ; no external lib
State / Sync	React Context + Firestore SDK	Real-time bi-directional sync
Routing	React Router v6	Client-side HTML5 history routing

Table 4.1: Front-end technology stack

4.2 Client-Side Neural Network (MLP)

The procrastination risk model is a **Multi-Layer Perceptron** written in pure JavaScript and executed entirely within the browser runtime. No training data or model weights ever leave the device.

4.2.1 Model Specification

Property	Value
Algorithm	Online supervised backpropagation
Training epochs	10 per update cycle
Output range	$p \in [0, 1]$ (procrastination probability)
Storage	localStorage (device-only)
Input Features	
Feature 1	Task urgency (normalised 0–1)
Feature 2	Estimated duration (hours)
Feature 3	Historical deferral rate
Feature 4	Time-of-day coefficient
Feature 5	Deadline proximity (hours remaining)

Table 4.2: MLP neural network specification

All neural network weights and training samples are stored exclusively in the user’s browser `localStorage`. No behavioural telemetry is transmitted to any external endpoint.

4.3 Firestore Security Rules

All Firestore reads and writes are protected by production-grade security rules that enforce strict authenticated owner-only access:

```
rules_version = '2';
service cloud.firestore {
  match /databases/{database}/documents {
    match /users/{userId}/{document=**} {
      allow read, write:
        if request.auth != null
          && request.auth.uid == userId;
    }
  }
}
```

These rules guarantee that:

- Unauthenticated requests are rejected at the database layer.
- Each user can only access documents within their own `/users/{userId}/` namespace.
- Cross-user data leakage is architecturally impossible.

4.4 Project Directory Structure

```
deadlineiq/  
  public/                # Static assets (favicon, icons)  
  src/  
    components/  
      Dashboard/        # HUD, telemetry, navigation  
      Calendar/         # Glassmorphic weekly grid  
      Kanban/           # Drag-and-drop progress board  
      Forensics/        # Procrastination analytics  
      Coach/            # IQ Coach co-pilot sidebar  
      Habits/           # Streak tracker + gauges  
      Settings/         # API keys and preferences  
    utils/  
      localML.js        # Client-side MLP neural network  
    services/  
      firebase.js       # Firestore and Auth SDK  
      gemini.js         # Gemini API client wrapper  
      calendar.js       # Google Calendar API wrapper  
    App.jsx             # Root component and routing  
    main.jsx            # Vite entry point  
  .env                  # Environment variables (NOT  
    committed)  
  firestore.rules       # Firestore security rules  
  firebase.json         # Firebase Hosting configuration  
  Dockerfile            # Multi-stage container build  
  nginx.conf            # SPA routing fallback config  
  vite.config.js        # Vite build configuration  
  package.json          # Dependencies and scripts
```


CHAPTER 5

Local Development Setup

5.1 Prerequisites

Tool	Required Version	Notes
Node.js	v18 or v20	LTS releases strongly recommended
npm	Bundled with Node	Alternatively use pnpm or yarn
Git	Any recent release	For cloning the repository

Table 5.1: Development prerequisites

5.2 Installation Steps

5.2.1 Step 1 Navigate and Install Dependencies

```
1 # Navigate to the project directory
2 cd deadlineiq
3
4 # Install all npm dependencies
5 npm install
```

5.2.2 Step 2 Environment Configuration

Create a `.env` file in the project root. Populate it with your Firebase project credentials:

```
VITE_FIREBASE_API_KEY=your_firebase_api_key
VITE_FIREBASE_AUTH_DOMAIN=your_project_id.firebaseio.com
VITE_FIREBASE_PROJECT_ID=your_project_id
VITE_FIREBASE_STORAGE_BUCKET=your_project_id.appspot.com
VITE_FIREBASE_MESSAGING_SENDER_ID=your_messaging_sender_id
VITE_FIREBASE_APP_ID=your_firebase_app_id
```

The Gemini Developer API Key can be entered directly in the web UI under **Settings** → **API Configuration**, so you do not need to add it to the `.env` file.

5.2.3 Step 3 Start the Development Server

```
1 npm run dev
```

Open <http://localhost:5174> in your browser. Vite's Hot Module Replacement (HMR) provides instant UI feedback during development.

CHAPTER 6

Deployment

6.1 Docker Deployment

DeadlineIQ ships with a multi-stage `Dockerfile` that compiles static assets and serves them via a custom `nginx.conf` with HTML5 history-mode fallback routing.

6.1.1 Build the Container Image

```
1 docker build -t deadlineiq .
```

6.1.2 Run the Container

```
1 # Maps container port 80 to host port 8080
2 docker run -d -p 8080:80 deadlineiq
```

Access the application at <http://localhost:8080>.

The container exposes port **80** internally (nginx default). Adjust `-p HOST_PORT:80` to fit your environment. Ensure no other service is bound to the chosen host port.

6.2 Firebase Hosting Deployment

```
1 # 1. Authenticate with your Google account
2 npx firebase login
3
4 # 2. Build optimised production assets
5 npm run build
6
7 # 3. Deploy to Firebase Hosting
8 npx firebase deploy --only hosting
```

After a successful deployment, the application is globally available at:

- `https://your-project-id.web.app`
- `https://your-project-id.firebaseio.com`

Both endpoints are served over verified SSL with Firebase's global CDN.

6.3 Google Cloud Run Deployment

For serverless, auto-scaling production hosting, the project includes pre-configured Cloud Run artifacts. Integrate `cloudbuild.yaml` into your CI/CD pipeline to trigger automated container builds and deployments on every push to the main branch.

Security Considerations

7.1 Authentication Layer

Firebase Authentication with Google OAuth 2.0 ensures every session is cryptographically signed and token-validated. Token expiry and refresh cycles are handled automatically by the Firebase SDK.

7.2 Data Isolation at the Database Layer

The Firestore namespace structure (`/users/{userId}/...`) combined with the security rules described in Section 4.3 makes cross-user data access architecturally impossible – not merely a policy decision.

7.3 Client-Side ML Privacy

The MLP neural network runs entirely within the browser’s JavaScript engine. Behavioural training data and learned weights are persisted only to `localStorage` on the user’s device.

7.4 API Key Security Best Practices

- Add `.env` to `.gitignore` immediately. **Never commit credentials to version control.**
- In the Google Cloud Console, restrict Firebase API keys to authorised domains only (e.g., your `web.app` domain).
- Rotate Gemini API keys periodically from the Google AI Studio dashboard.
- Enable Firebase App Check to block unauthorised clients from accessing your backend.
- Review Firestore security rules before every production release.

Technical Documentation • v1.0 • June 2026

For configuration help, contribution guidelines, or bug reports,
consult the repository `README.md` or contact the engineering team.
